Classification metrics optimization: Logloss and accuracy

Classification metrics optimization

- Logloss

- Accuracy

- AUC

- (Quadratic weighted) Kappa

Logloss

$$\text{LogLoss} = -\frac{1}{N}\sum_{i=1}^{N} y_i \log(\hat{y}_i) + (1 - y_i)\log(1 - \hat{y}_i)$$

How do you optimize it?

Just run the right model!
(or calibrate others)

Logloss

- **Tree-based**

  XGBoost, LightGBM

  ~~sklearn.RandomForestClassifier~~

- **Linear models**

  sklearn.<>Regression

  sklearn.SGDRegressor

  Vowpal Wabbit

- **Neural nets**

  PyTorch, Keras, TF, etc.

Synonyms: Logistic loss   Read the docs!
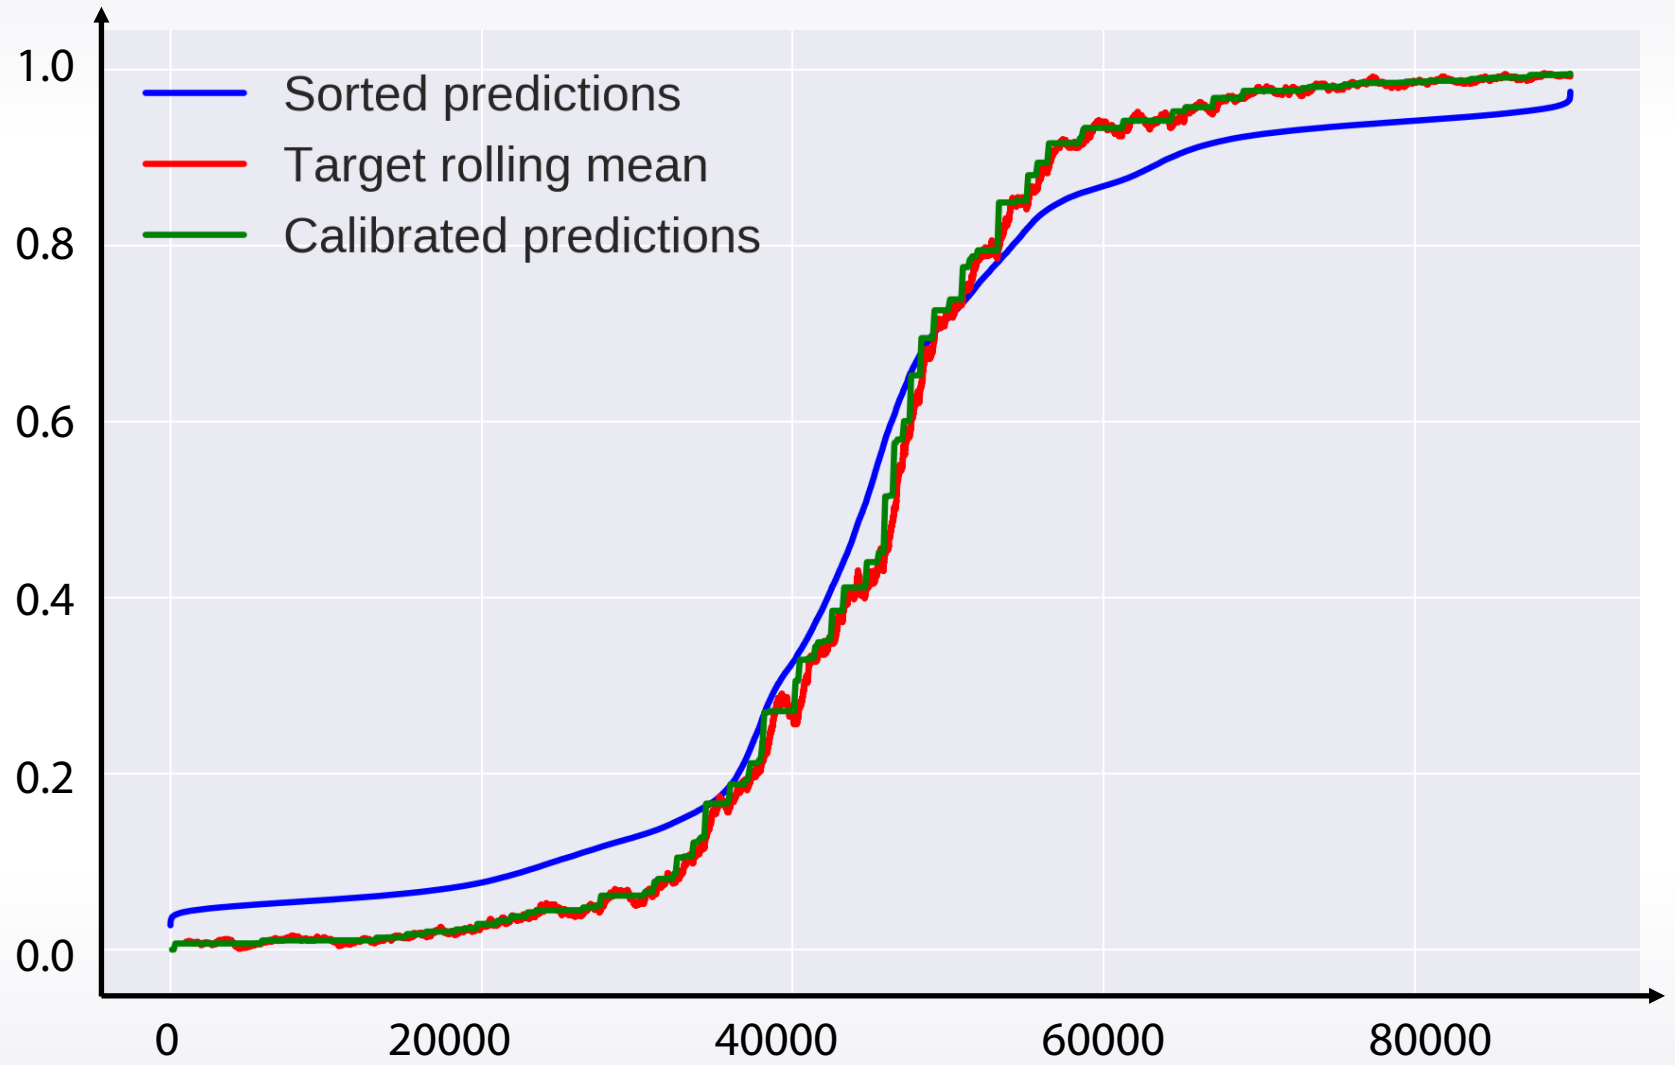
Logloss

Correct probabilities:

- Take all objects with score e.g. ~ 0.8
  - 80% of them of class 1
  - 20% of them class 0

Incorrect probabilities:

- Take all objects with score e.g. ~ 0.8
  - 50% of them of class 1
  - 50% of them of class 0

Probability calibration

Probability calibration

- Platt scaling
  - Just fit Logistic Regression to your predictions (like in stacking)

- Isotonic regression
  - Just fit Isotonic Regression to your predictions (like in stacking)

- Stacking
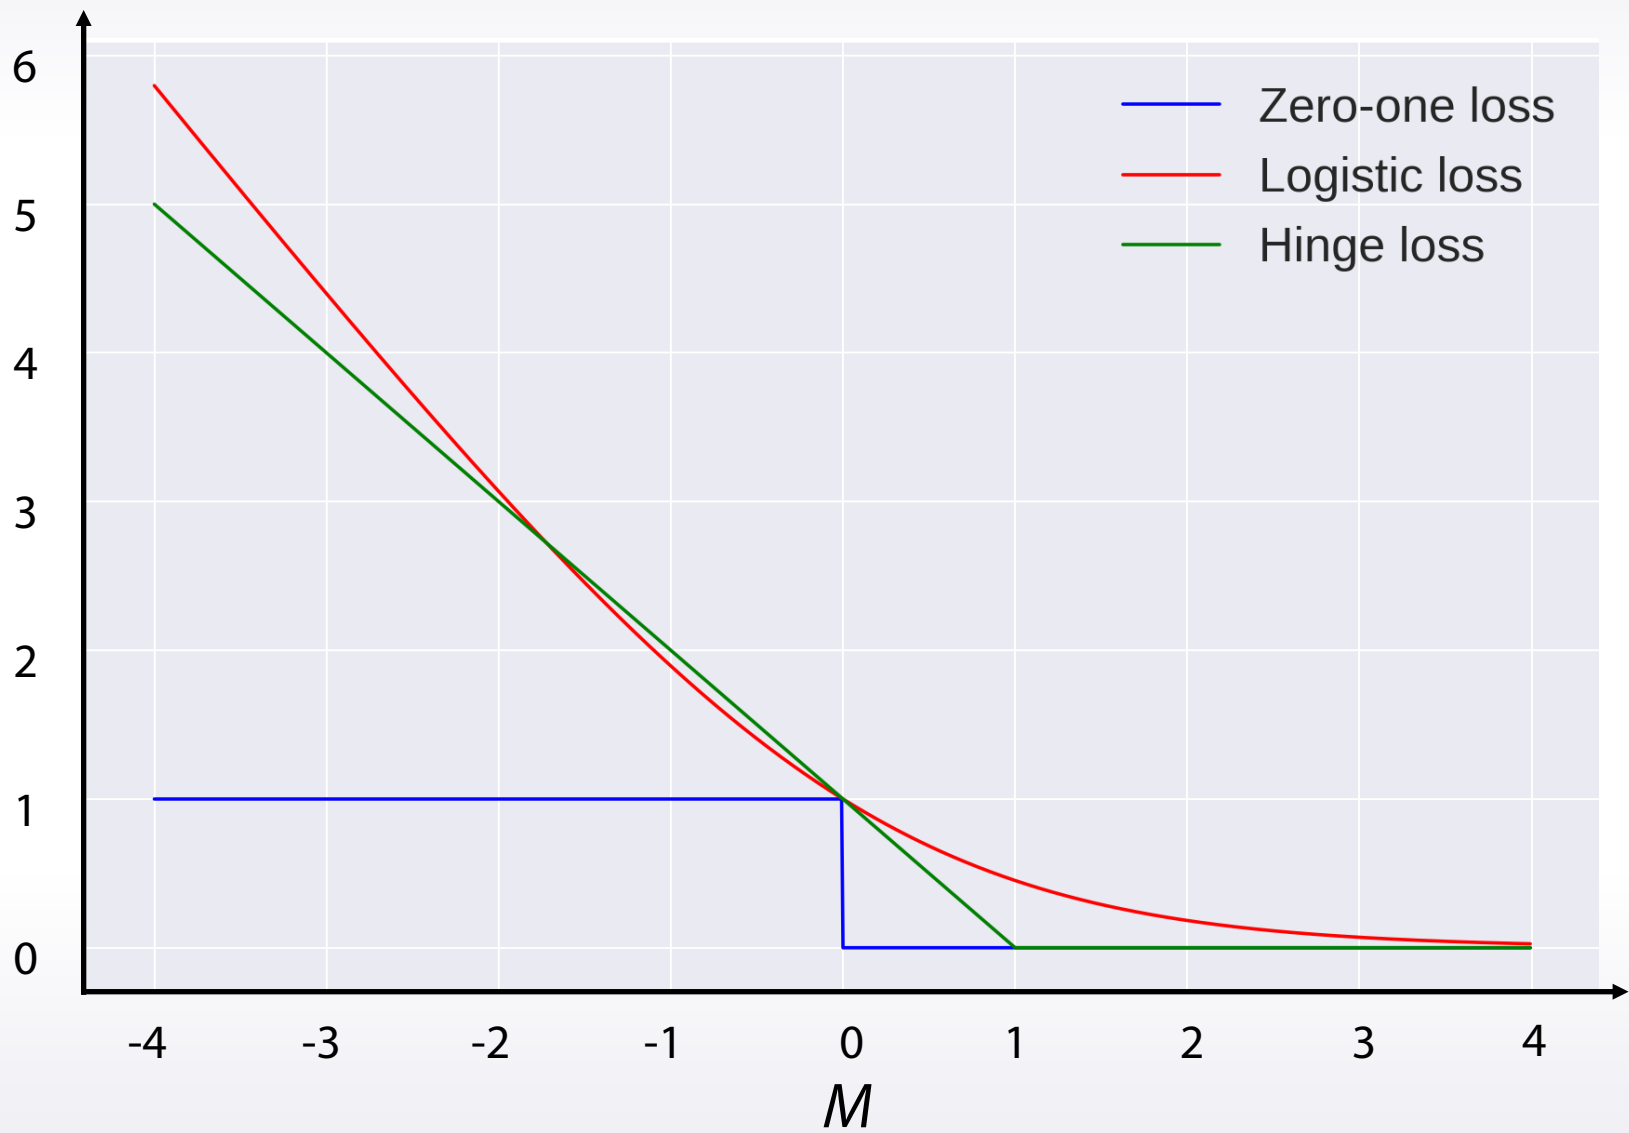 - Just fit XGBoost or neural net to your predictions

Accuracy

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^{N} [\hat{y}_i = y_i]$$
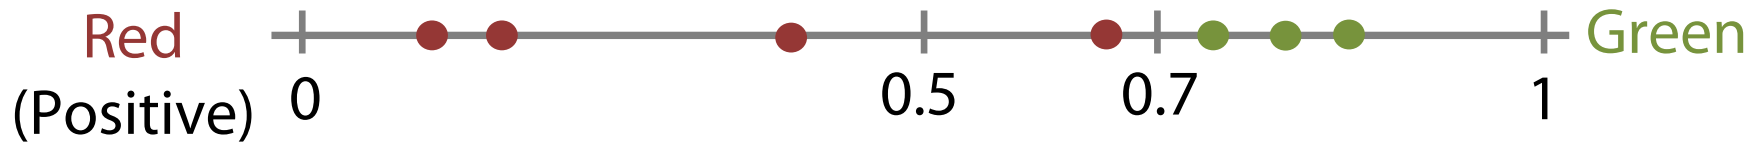
How do you optimize it?

Fit any metric and tune treshold!

Accuracy

Accuracy

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^{N} [[f(x) > b] = y_i]$$

Red (Positive) — Green

$$b = 0.5 \quad \Rightarrow \quad \text{Accuracy} = \frac{6}{7}$$

$$b = 0.7 \quad \Rightarrow \quad \text{Accuracy} = 1$$

Conclusion

- Logloss

- Accuracy

- AUC

- (Quadratic weighted) Kappa